Nityashree S

+91 78270 76283 · nitya08shree@gmail.com · **Portfolio** · LinkedIn · GitHub · India

AI engineer who builds agents that hold up in production: safe with people's data, tested before release, and worth what they cost to run. I came to AI by testing an assistant before launch, and I still work that way. I direct AI coding agents and prove every change before it ships.

EXPERIENCE

AI Engineer

Jun 2026 – Present

AI Team · Yuno · Onsite

PLATFORM AND RELIABILITY

- Found why a production agent went offline (one badly formatted instruction file), proved it with a failing test, and fixed it so a single bad file can no longer take an agent down.
- Added time limits, metrics and clear logs to the gateway every agent tool call passes through. The new logs exposed a connector that had been failing silently, which I then fixed.

ACCESS AND SECURITY

- Red-teamed the new agent engine before launch: chained a leaked system prompt into removing the agent's safety rules and poisoning its memory, and filed 25 bug reports. The fix moved those rules out of the prompt and into the platform.
- Built guest access so people can safely share their AI assistant with teammates, across five services: each guest gets their own permissions and their own conversation.
- Built a connector where each employee signs in with their own account, so their agent can only see what they are allowed to see. Live in production.
- Closed security gaps: a link that could log users out from another website, chat events accepted without proof of origin, and checks that let requests through when a secret was missing (they now refuse).

EVALS AND MODEL BEHAVIOUR

- Built the step that decides which language each agent reply uses, and tested it silently beside the old logic on a labelled test set before release.
- Showed with a small A/B test (same prompt, different model) that the model, not the prompt, made agents ignore instructions; the team changed the default model.

AGENTS FOR BUSINESS TEAMS

- Built finance agents: invoice intake, where the model reads and tested code decides, and a mailbox sorter whose projected model cost I cut by changing its schedule after measuring empty runs.
- Worked day to day with teammates and stakeholders across India and Latin America.

Letter of recommendation from Simón Martínez, VP of AI at Yuno, available on request.

Payments Analyst Intern

Dec 2025 – Jun 2026

Office of the CEO · Yuno · Onsite

- Built and maintained a database of which payment providers and methods work in which markets, and audited it for errors (duplicates, wrong categories, wrong regions).
- Tested an internal AI assistant before launch with real questions from Sales and Business Development, and caught a data error behind a wrong answer.
- Built an interactive KPI dashboard for team leads by directing AI coding tools; it pulls live data from Salesforce, MongoDB and Google Sheets through their APIs.
- Automated manual work across the company, above all keeping partnership and payment-provider information up to date, so teams stopped updating it by hand.
- Worked directly with leadership, Sales and Partnerships to turn their questions into data they could act on.

PROJECTS

- **Smart Router** · LLM gateway · Python, FastAPI, Redis, Postgres, Docker
  Sends each prompt to a cheaper or stronger AI model by simple rules, falls back if a provider fails, and logs the estimated cost of every request.

- **Meridian** · multi-agent research · Python, FastAPI, SQLite
  A planner agent splits a question, researcher agents search in parallel, a critic agent flags gaps, and every step streams live.
- **Cross-Border Fee API** · payments · TypeScript, Express, Jest · Live demo
  An itemised price for a Latin American payment: currency spread, provider fees and local taxes.

SKILLS

AI agents	LLM agents and tool calling, prompt engineering, MCP connectors, multi-agent workflows, scheduled agents
Evals and QA	LLM evals with labelled test sets, shadow testing, A/B model comparison, acceptance testing, root-cause analysis
AI security	Red-teaming, prompt-injection defence, access control, OAuth sign-in, webhook verification
Production	Logging and metrics, incident debugging, feature-flagged releases, LLM cost optimisation
Stack	Python, TypeScript, Go, FastAPI, Next.js, Postgres, MongoDB, Redis, Docker, Vercel, Salesforce and Google Sheets APIs (built by directing AI coding agents; I review, test and own it)
Tools	Claude Code, Claude Design, Cursor, DeepSeek, Grok; personal agents in Slack for my pull-request workflow
Ways of working	Gathering requirements from business teams, payments domain

EDUCATION

B.Tech, Computer Science and Engineering (AI & Data Science) · SASTRA University 2022 – 2026

Class XII, CBSE · Convent of Jesus and Mary

CERTIFICATIONS

- **OWASP Top 10:2025 in Go** · SecureFlag · 2026
- **Introduction to Generative AI** and **Prompt Design in Vertex AI** (skill badge) · Google Cloud · 2024
- **Supervised Machine Learning: Regression and Classification** · DeepLearning.AI and Stanford University (Coursera) · 2024
- **Prompt Engineering for Everyone** and **Create a Voice Assistant with OpenAI GPT-3 and IBM Watson** · IBM Cognitive Class · 2024
- **Artificial Intelligence Use Cases and Applications** · AWS Training and Certification · 2024
- **Machine Learning Pipelines with Azure ML Studio** · Coursera · 2024

LANGUAGES

English (full professional proficiency) · Hindi (native) · Tamil (native)